

容器与算法

主讲教师： 沈艳艳

Vector: Changing Size

Changing vector size

- Fundamental problem addressed
 - We (humans) want abstractions that can change size (*e.g.*, a vector where we can change the number of elements). However, in computer memory everything must have a fixed size, so how do we create the illusion of change?

- Given

vector v(n); // v.size()==n

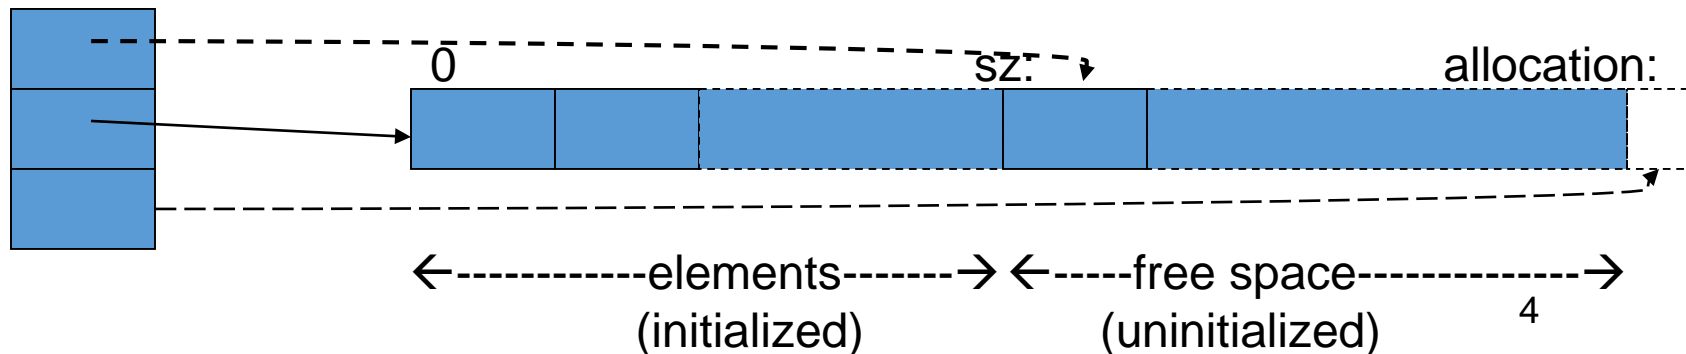
we can change its size in three ways

- Resize it
 - **v.resize(10); // v now has 10 elements**
- Add an element
 - **v.push_back(7); // add an element with the value 7 to the end of v**
// v.size() increases by 1
- Assign to it
 - **v = v2; // v is now a copy of v2**
// v.size() now equals v2.size()

Representing vector


- If you **resize()** or **push_back()** once, you'll probably do it again;
 - let's prepare for that by sometimes keeping a bit of free space for future expansion

```
class vector {  
    int sz;  
    double* elem;  
    int space; // number of elements plus "free space"  
               // (the number of "slots" for new elements)  
  
public:  
    // ...  
};
```




Representing vector

- An empty vector (no free store use):



- A vector(n) (no free space):



vector::reserve()

- First deal with space (allocation); given space all else is easy
 - Note: **reserve()** doesn't mess with size or element values

void vector::reserve(int newalloc)

*// make the vector have space for **newalloc** elements*

{

if (newalloc<=space) return; *// never decrease allocation*

double* p = new double[newalloc]; *// allocate new space*

for (int i=0; i<sz; ++i) p[i]=elem[i]; *// copy old elements*

delete[] elem; *// deallocate old space*

elem = p;

space = newalloc;

}

vector::resize()

- Given **reserve()**, **resize()** is easy
 - **reserve()** deals with space/allocation
 - **resize()** deals with element values

void vector::resize(int newsize)

*// make the vector have **newsize** elements*

// initialize each new element with the default value 0.0

{

reserve(newsize); *// make sure we have sufficient space*

for(int i = sz; i<newsize; ++i) elem[i] = 0; *// initialize new elements*

sz = newsize;

}

vector::push_back()

- Given **reserve()**, **push_back()** is easy
 - **reserve()** deals with space/allocation
 - **push_back()** just adds a value

```
void vector::push_back(double d)
```

```
    // increase vector size by one
```

```
    // initialize the new element with d
```

```
{
```

```
    if (sz==0)                // no space: grab some
```

```
        reserve(8);
```

```
    else if (sz==space)       // no more free space: get more space
```

```
        reserve(2*space);
```

```
    elem[sz] = d;             // add d at end
```

```
    ++sz;                     // and increase the size (sz is the number of elements)
```

```
}
```


resize() and push_back()

```
class vector { // an almost real vector of doubles
    int sz; // the size
    double* elem; // a pointer to the elements
    int space; // size+free_space
public:
    // ... constructors and destructors ...

    double& operator[ ](int n) { return elem[n]; } // access: return reference
    int size() const { return sz; } // current size

    void resize(int newsize); // grow
    void push_back(double d); // add element

    void reserve(int newalloc); // get more space
    int capacity() const { return space; } // current available space
};
```

Assignment

- Copy and swap is a powerful general idea

vector& vector::operator=(const vector& a)

// like copy constructor, but we must deal with old elements

*// make a copy of **a** then replace the current **sz** and **elem** with **a**'s*

{

double* p = new double[a.sz]; *// allocate new space*

for (int i = 0; i < a.sz; ++i) p[i] = a.elem[i]; *// copy elements*

delete[] elem; *// deallocate old space*

sz = a.sz; *// set new size*


elem = p; *// set new elements*

return *this; *// return a self-reference*

}

Optimize assignment

- “Copy and swap” is the most general idea
 - but not always the most efficient
 - What if there already is sufficient space in the target vector?
 - Then just copy!
 - For example: **a = b;**



Optimized assignment

```
vector& vector::operator=(const vector& a)
{
    if (this==&a) return *this;    // self-assignment, no work needed

    if (a.sz<=space) {             // enough space, no need for new allocation
        for (int i = 0; i<a.sz; ++i) elem[i] = a.elem[i];    // copy elements
        space += sz-a.sz;           // increase free space
        sz = a.sz;
        return *this;
    }

    double* p = new double[a.sz];    // copy and swap
    for (int i = 0; i<a.sz; ++i) p[i] = a.elem[i];
    delete[ ] elem;
    sz = a.sz;
    space = a.sz;
    elem = p;
    return *this;
}
```

Templates and Range Checking

Templates

- But we don't just want vector of double
- We want vectors with element types we specify
 - `vector<double>`
 - `vector<int>`
 - `vector<Month>`
 - `vector<Record*>` *// vector of pointers*
 - `vector<vector<Record>>` *// vector of vectors*
 - `vector<char>`
- We must make the element type a parameter to **vector**
- **vector** must be able to take both built-in types and user-defined types as element types
- This is not some magic reserved for the compiler; we can define our own parameterized types, called “templates”

Templates

- The basis for generic programming in C++
 - Sometimes called “parametric polymorphism”
 - Parameterization of types (and functions) by types (and integers)
 - Unsurpassed flexibility and performance
 - Used where performance is essential (e.g., hard real time and numerics)
 - Used where flexibility is essential (e.g., the C++ standard library)

- Template definitions

```
template<class T, int N> class Buffer { /* ... */ };  
template<class T, int N> void fill(Buffer<T,N>& b) { /* ... */ }
```

- Template specializations (instantiations)

// for a class template, you specify the template arguments:

```
Buffer<char,1024> buf;           // for buf, T is char and N is 1024
```

// for a function template, the compiler deduces the template arguments:

```
fill(buf); // for fill(), T is char and N is 1024; that's what buf has
```

Parameterize with element type

*// an almost real **vector** of **Ts**:*

```
template<class T> class vector {
```

```
    // ...
```

```
};
```

```
vector<double> vd;           // T is double
```

```
vector<int> vi;             // T is int
```

```
vector<vector<int>> vvi;    // T is vector<int>
```

```
// in which T is int
```

```
vector<char> vc;           // T is char
```

```
vector<double*> vpd;        // T is double*
```

```
vector<vector<double>*> vvpd; // T is vector<double>*
```

```
// in which T is double
```


Basically, `vector<double>` is

*// an almost real **vector** of **doubles**:*

```
class vector {  
    int sz;           // the size  
    double* elem;     // a pointer to the elements  
    int space;        // size+free_space  
public:  
    vector() : sz(0), elem(0), space(0) { }           // default constructor  
    vector(const vector&);                             // copy constructor  
    vector& operator=(const vector&);                 // copy assignment  
    ~vector() { delete[ ] elem; }                     // destructor  
  
    double& operator[ ] (int n) { return elem[n]; }   // access: return  
                                                    reference  
    int size() const { return sz; }                   // the current size  
  
    // ...  
};
```

Basically, **vector<char>** is

*// an almost real **vector** of **chars**:*

```
class vector {  
    int sz;           // the size  
    char* elem;      // a pointer to the elements  
    int space;        // size+free_space  
public:  
    vector() : sz{0}, elem{0}, space{0} { }           // default constructor  
    vector(const vector&);                             // copy constructor  
    vector& operator=(const vector&);                 // copy assignment  
    ~vector() { delete[ ] elem; }                    // destructor  
  
    char& operator[ ] (int n) { return elem[n]; }      // access: return reference  
    int size() const { return sz; }                   // the current size  
  
    // ...  
};
```

Basically, **vector<T>** is

*// an almost real **vector** of Ts:*

template<class T> class vector { *// read “for all types T” (just like in math)*

int sz; *// the size*

T* elem; *// a pointer to the elements*

int space; *// size+free_space*

public:

vector() : sz{0}, elem{0}, space{0}; *// default constructor*

vector(const vector&); *// copy constructor*

vector& operator=(const vector&); *// copy assignment*

vector(const vector&&); *// move constructor*

vector& operator=(vector&&); *// move assignment*

~vector() { delete[] elem; } *// destructor*

// ...

};

Basically, **vector<T>** is

*// an almost real **vector** of Ts:*

template<class T> class vector { *// read “for all types T” (just like in math)*

int sz; *// the size*

T* elem; *// a pointer to the elements*

int space; *// size+free_space*

public:

// ... constructors and destructors ...

T& operator[] (int n) { return elem[n]; }

// access: return reference

int size() const { return sz; }

// the current size

void resize(int newsize);

// grow

void push_back(double d);

// add element

void reserve(int newalloc);

// get more space

int capacity() const { return space; }

// current available space

// ...

};

Templates

- Problems (“there is no free lunch”)
 - Poor error diagnostics
 - Often spectacularly poor (but getting better in C++11; much better in C++14)
 - Delayed error messages
 - Often at link time
 - All templates must be fully defined in each translation unit
 - So place template definitions in header files
- Recommendation
 - Use template-based libraries
 - Such as the C++ standard library
 - *E.g.*, **vector**, **sort()**
 - Soon to be described in some detail
 - Initially, write only very simple templates yourself
 - Until you get more experience

Range checking

*// an almost real **vector** of Ts:*

```
struct out_of_range { /* ... */ };
```

```
template<class T> class vector {  
    // ...  
    T& operator[ ](int n);           // access  
    // ...  
};
```

```
template<class T> T& vector<T>::operator[ ](int n)  
{  
    if (n<0 || sz<=n) throw out_of_range();  
    return elem[n];  
}
```

Range checking

```
void fill_vec(vector<int>& v, int n)           // initialize v with factorials
{
    for (int i=0; i<n; ++i) v.push_back(factorial(i));
}

int main()
{
    vector<int> v;
    try {
        fill_vec(v,10);
        for (int i=0; i<=v.size(); ++i)
            cout << "v[" << i << "]==" << v[i] << '\n';
    }
    catch (out_of_range) {                    // we'll get here (why?)
        cout << "out of range error";
        return 1;
    }
}
```

Exception handling

- We use exceptions to report errors
- We must ensure that use of exceptions
 - Doesn't introduce new sources of errors
 - Doesn't complicate our code
 - Doesn't lead to resource leaks

STL

(The containers, iterators and algorithms)

- STL – the containers and algorithms part of the C++ standard library

Common tasks

- Collect data into containers
- Organize data
 - For printing
 - For fast access
- Retrieve data items
 - By index (e.g., get the **N**th element)
 - By value (e.g., get the first element with the value "**Chocolate**")
 - By properties (e.g., get the first elements where "**age<64**")
- Add data
- Remove data
- Sorting and searching
- Simple numeric operations

Observation

We can (already) write programs that are very similar independent of the data type used

- Using an **int** isn't that different from using a **double**
- Using a **vector<int>** isn't that different from using a **vector<string>**

Ideals

We'd like to write common programming tasks so that we don't have to re-do the work each time we find a new way of storing the data or a slightly different way of interpreting the data

- Finding a value in a **vector** isn't all that different from finding a value in a **list** or an array
- Looking for a **string** ignoring case isn't all that different from looking at a **string** not ignoring case
- Graphing experimental data with exact values isn't all that different from graphing data with rounded values
- Copying a file isn't all that different from copying a vector

Ideals (continued)

- Code that's
 - Easy to read
 - Easy to modify
 - Regular
 - Short
 - Fast
- Uniform access to data
 - Independently of how it is stored
 - Independently of its type
- ...

Ideals (continued)

- ...
- Type-safe access to data
- Easy traversal of data
- Compact storage of data
- Fast
 - Retrieval of data
 - Addition of data
 - Deletion of data
- Standard versions of the most common algorithms
 - Copy, find, search, sort, sum, ...


Examples

- Sort a vector of strings
- Find an number in a phone book, given a name
- Find the highest temperature
- Find all values larger than 800
- Find the first occurrence of the value 17
- Sort the telemetry records by unit number
- Sort the telemetry records by time stamp
- Find the first value larger than “Petersen”?
- What is the largest amount seen?
- Find the first difference between two sequences
- Compute the pairwise product of the elements of two sequences
- What are the highest temperatures for each day in a month?
- What are the top 10 best-sellers?
- What’s the entry for “C++” (say, in Google)?
- What’s the sum of the elements?

The STL

- Part of the ISO C++ Standard Library
- Mostly non-numerical
 - Only 4 standard algorithms specifically do computation
 - Accumulate, inner_product, partial_sum, adjacent_difference
 - Handles textual data as well as numeric data
 - E.g. string
 - Deals with organization of code and data
 - Built-in types, user-defined types, and data structures
- Optimizing disk access was among its original uses
 - Performance was always a key concern

The STL




- Designed by Alex Stepanov
- General aim: The most general, most efficient, most flexible representation of concepts (ideas, algorithms)
 - Represent separate concepts separately in code
 - Combine concepts freely wherever meaningful
- General aim to make programming “like math”
 - or even “Good programming *is* math”
 - works for integers, for floating-point numbers, for polynomials, for ...

Basic model

- Algorithms

sort, find, search, copy, ...



- Containers

vector, list, map, unordered_map, ...

- Separation of concerns


- Algorithms manipulate data, but don't know about containers
- Containers store data, but don't know about algorithms
- Algorithms and containers interact through iterators
 - Each container has its own iterator types

The STL

- An ISO C++ standard framework of about 10 containers and about 60 algorithms connected by iterators
 - Other organizations provide more containers and algorithms in the style of the STL
 - Boost.org, Microsoft, SGI, ...
- Probably the currently best known and most widely used example of generic programming

Basic model

- A pair of iterators defines a sequence
 - The beginning (points to the first element – if any)
 - The end (points to the one-beyond-the-last element)




- An iterator is a type that supports the “iterator operations”
 - ++ Go to next element
 - * Get value
 - == Does this iterator point to the same element as that iterator?
- Some iterators support more operations (e.g. --, +, and [])

Containers


(hold sequences in difference ways)

- **vector**




- **list**

(doubly linked)




- **set**

(a kind of tree)



The simplest algorithm: **find()**



// Find the first element that equals a value

```
template<class In, class T>
In find(In first, In last, const T& val)
{
    while (first!=last && *first != val) ++first;
    return first;
}
```

```
void f(vector<int>& v, int x) // find an int in a vector
{
    vector<int>::iterator p = find(v.begin(),v.end(),x);
    if (p!=v.end()) { /* we found x */ }
    // ...
}
```

We can ignore (“abstract away”) the differences between containers

find()

generic for both element type and container type


```
void f(vector<int>& v, int x)           // works for vector of ints
{
    vector<int>::iterator p = find(v.begin(),v.end(),x);
    if (p!=v.end()) { /* we found x */ }
    // ...
}
```

```
void f(list<string>& v, string x)       // works for list of strings
{
    list<string>::iterator p = find(v.begin(),v.end(),x);
    if (p!=v.end()) { /* we found x */ }
    // ...
}
```


```
void f(set<double>& v, double x)       // works for set of doubles
{
    set<double>::iterator p = find(v.begin(),v.end(),x);
    if (p!=v.end()) { /* we found x */ }
    // ...
}
```

Algorithms and iterators

- An iterator points to (refers to, denotes) an element of a sequence
- The end of the sequence is “one past the last element”
 - **not** “the last element”
 - That’s necessary to elegantly represent an empty sequence
 - One-past-the-last-element isn’t an element
 - You can compare an iterator pointing to it
 - You can’t dereference it (read its value)
- Returning the end of the sequence is the standard idiom for “not found” or “unsuccessful”



An empty sequence:




Simple algorithm: `find_if()`

- Find the first element that matches a criterion (predicate)
 - Here, a predicate takes one argument and returns a **bool**

```
template<class In, class Pred>
In find_if(In first, In last, Pred pred)
{
    while (first!=last && !pred(*first)) ++first;
    return first;
}
```

```
void f(vector<int>& v)
{
    vector<int>::iterator p = find_if(v.begin(),v.end,Odd());
    if (p!=v.end()) { /* we found an odd number */ }
    // ...
}
```

A predicate



Predicates

- A predicate (of one argument) is a function or a function object that takes an argument and returns a **bool**
- For example

- A function

```
bool odd(int i) { return i%2; } // % is the remainder (modulo)
    operator
odd(7);                        // call odd: is 7 odd?
```

- A function object

```
struct Odd {
    bool operator()(int i) const { return i%2; }
};
Odd odd;    // make an object odd of type Odd
odd(7);     // call odd: is 7 odd?
```

vector

```
template<class T> class vector {  
    T* elements;  
    // ...  
    using value_type = T;  
    using iterator = ???;           // the type of an iterator is implementation defined  
                                     // and it (usefully) varies (e.g. range checked iterators)  
                                     // a vector iterator could be a pointer to an element  
    using const_iterator = ???;  
  
    iterator begin();                // points to first element  
    const_iterator begin() const;  
    iterator end();                  // points to one beyond the last element  
    const_iterator end() const;  
  
    iterator erase(iterator p);       // remove element pointed to by p  
    iterator insert(iterator p, const T& v); // insert a new element v before p  
};
```

Link:

T value

Link* pre
Link* post

list

```
template<class T> class list {  
    Link* elements;  
    // ...  
    using value_type = T;  
    using iterator = ???;           // the type of an iterator is implementation defined  
                                     // and it (usefully) varies (e.g. range checked iterators)  
                                     // a list iterator could be a pointer to a link node  
    using const_iterator = ???;  
  
    iterator begin();                // points to first element  
    const_iterator begin() const;  
    iterator end();                  // points to one beyond the last element  
    const_iterator end() const;  
  
    iterator erase(iterator p);       // remove element pointed to by p  
    iterator insert(iterator p, const T& v); // insert a new element v before p  
};
```

Vector vs. List

- By default, use a **vector**
 - You need a reason not to
 - You can “grow” a vector (e.g., using **push_back()**)
 - You can **insert()** and **erase()** in a vector
 - Vector elements are compactly stored and contiguous
 - For small vectors of small elements all operations are fast
 - compared to lists
- If you don’t want elements to move, use a **list**
 - You can “grow” a list (e.g., using **push_back()** and **push_front()**)
 - You can **insert()** and **erase()** in a list
 - List elements are separately allocated
- Note that there are more containers, e.g.,
 - map
 - unordered_map

Some useful standard headers

- **<iostream>** I/O streams, cout, cin, ...
- **<fstream>** file streams
- **<algorithm>** sort, copy, ...
- **<numeric>** accumulate, inner_product, ...
- **<functional>** function objects
- **<string>**
- **<vector>**
- **<map>**
- **<unordered_map>** hash table
- **<list>**
- **<set>**